

mmap

crazy big, crazy fast, crazy flexible

Presented July 25, 2018
CRUG + ChiPy, Chicago IL, USA

Jeffrey A. Ryan jeff.a.ryan@gmail.com

DISCLAIMER

These are my **opinions**

Only my **opinions**

Not my employer's opinions

Not friend's opinions (especially not Dirk's :)

No warranty or guarantees

But *maybe* this will be useful.

Maybe you've seen this already?

mmap + indexing

crazy big data, crazy fast

Presented May 27, 2010
Inaugural Chicago RUG

Jeffrey A. Ryan
insight algorithmics, inc

insight algorithmics, inc
Jeffrey A. Ryan

Inaugural Chicago RUG
Presented May 27, 2010

Background

R and Python since 2003

Co-founder of R/Finance Conference 2009

Quant in Hedge Fund

Created quantmod, xts, etc.

DataCamp course on xts

What is mmap ?

memory-mapped file (like)

system call (core operating system)

constrained by address space, not memory

fast access - once “mapped” behaves like RAM

Why does this matter ?


larger than memory datasets

language agnostic - just bytes

cross language sharing is easy (R + Python)

Isn't this like a database ? YES!

Database Design 101



Row or Column Based


By Row



By Column



R is Column Oriented



i.e.

data.frames store column
values sequentially

a.k.a column-major

column 1




...

column 2



...

column 3



Why mmap ?

R is memory limited

Need memory many times data size

Access isn't always columnar

Maybe you access all row values together ?

Sharing

Pipeline requires different tools - language agnostic serialization!

Random access

Data access isn't always linear

mmap package

OS system call

very low level API - you see what the C call sees

virtually map files into memory on demand

mapped files are treated *as if* they are memory

cross-platform - Windows and UNIX-alikes

mmap similar (but different) to the R packages *ff* and *bigmemory*

mmap package

mmap	R	C	bytes
raw()	raw	unsigned char	1
bits()	integer	int	1/32
char()	raw	char	1
uchar()	raw	unsigned char	1
int8()	integer	signed char	1
uint8()	integer	unsigned char	1
int16()	integer	signed short	2
uint16()	integer	unsigned short	2
int24()	integer	three byte int	3
uint24()	integer	unsigned three byte int	3
int32()	integer	int	4
integer()	integer	int	4
real32()	double	single precision float	4
real64()	double	double precision float	8
double()	double	double precision float	8
cplx()	complex	complex	16
complex()	complex	complex	16
char(n)	character	fixed-width ascii	n+1
char(n,nul=F)	character	non-nul terminated	n
character(n)	character	fixed-width ascii	n+1
struct(...)	list	struct of above types	variable

Big Endian Support as of Summer 2018 !

mmap package


```
> # 4-byte (int32 or integer)
> # 8-byte (int64)
> # 8-byte float (real64 or double)
```

```
> m <- mmap(tmp, int32())
> m
<mmap:/var/folde...> (int) int [1:10000000] 1 2 3 4 5 6 ...
```

```
> m <- mmap(tmp, int64())
> m
<mmap:/var/folde...> (int64) num [1:10000000] 1 2 3 4 5 6 ...
> file.info(tmp)$size
[1] 8e+07
```

```
> m[34]
[1] 34
```

By Vector



mmap package

```
> # 2-byte (int16)
> # 4-byte (int32 or integer)
> # 8-byte float (real64 or double)

> record.type <- struct(short=int16(),int=int32(),double=real64())
> record.type

struct: (short) integer(0)
      (int) integer(0)
      (double) double(0)

> nbytes(record.type) # 14 bytes in total
[1] 14

> m <- mmap(tmp, record.type)
> m[1]
$short
[1] 1
$int
[1] 366214
$double
[1] -1.382365
```

By Row



mmap package

> example(mmap)

> example(struct)

> example(types)

Further Information

<http://cran.r-project.org/web/packages/mmap/index.html>

<http://past.rinfinance.com/agenda/2012/talk/JeffRyan.pdf>

mmap + indexing

Real World Example

67,836,671 equity option contracts
13 columns, 12GB on disk

```
> system.time( db[symbols=="AAPL"] )  
   user  system elapsed  
0.012  0.000  0.012
```

```
> db[symbols=="AAPL"]  
91428 hits
```


mmap + indexing

Real World Example

6 queries in 3.3 seconds

get a single contract as an xts time-series given OSI key

last 3 days of all AAPL April calls that have a delta at some point between .5 and .8, showing bid,ask,iv, and volume as an xts time-series

number of records on April 13

osi, bid and ask of AAPL puts ($\text{delta} < 0$) on April 13, expiring on the April 17

same, sorted by decreasing iv, excluding no-bid contracts, limit to 15

plot 3 day EMA of bid-ask spread of AAPL options with IV between 20% and 30%